

MNIST

In [1]:

```
import numpy as np
import torch
import torch.utils.data as tud
import torch.nn as tnn
```

In [2]:

```
import matplotlib.pyplot as plt
```

In [3]:

```
from copy import deepcopy as dcopy
```

In [4]:

```
import torchvision as tv
import torchvision.datasets as tvds
```

In [5]:

```
mnist_trainset = tvds.MNIST(root='./data', train=True, download=True, transform=tv.transforms.ToTensor())
mnist_testset = tvds.MNIST(root='./data', train=False, download=True, transform=tv.transforms.ToTensor())
```

In [6]:

```
train_set, val_set = tud.random_split(mnist_trainset, [50000, 10000], generator=torch.Generator().manual_seed(0))
```

In [7]:

```
test_set = mnist_testset
```

In [8]:

```
del mnist_testset, mnist_trainset
```

In [9]:

```
# batch size increased for faster training
train_loader = tud.DataLoader(train_set, batch_size=64, shuffle=True, drop_last=True)
val_loader = tud.DataLoader(val_set, batch_size=10000, shuffle=True, drop_last=True)
test_loader = tud.DataLoader(test_set, batch_size=10000, shuffle=True, drop_last=True)
```

In [10]:

```
class my_nn_2(tnn.Module):
    def __init__(self, h=128, d_in=784, d_out=10):
        super(my_nn_2, self).__init__()
        self.linear1 = tnn.Linear(d_in, h)
        self.linear2 = tnn.Linear(h, d_out)
        # https://medium.com/@zhang_yang/understanding-cross-entropy-implementation-in-pytorch-h-softmax-log-softmax-nll-cross-entropy-416a2b200e34
        # softmax + nll loss is worse than cross entropy loss

    def forward(self, x):
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred
```

In [11]:

```
crit_bce = tnn.CrossEntropyLoss()
crit_mse = tnn.MSELoss()
```

In [12]:


```
def fit(h=128, epochs=20, criterion=tnn.CrossEntropyLoss(), lr=1e-3):
    model = my_nn_2(h)
    opt = torch.optim.SGD(model.parameters(), lr=lr)
    tlh=[]
    tah=[]
    vlh=[]
    vah=[]
    vl_min=10**10
    for ep in range(epochs):
        model.train()
        training_loss=0
        acc=0
        for (idx, b) in enumerate(train_loader):
            xb=b[0].reshape(-1,784).float()
            yb=b[1]
            yb_pred = model(xb)
            loss=criterion(yb_pred, yb)
            training_loss+=loss
            acc+=torch.sum(torch.argmax(yb_pred, axis=1)==yb)
            opt.zero_grad()
            loss.backward()
            opt.step()
        tah.append(acc/50000)
        tlh.append(training_loss/50000)
        model.eval()
        with torch.no_grad():
            val_loss=0
            acc=0
            for (idx, b) in enumerate(val_loader):
                xb=b[0].reshape(-1,784).float()
                yb=b[1]
                yb_pred = model(xb)
                loss=criterion(yb_pred, yb)
                acc+=torch.sum(torch.argmax(yb_pred, axis=1)==yb)
                val_loss+=loss
            vah.append(acc/10000)
            vl=val_loss/64
            if vl<vl_min:
                best_model=dcopy(model)
                vl_min=vl
            vlh.append(vl)
    return best_model, tlh, tah, vlh, vah
```


In [13]:

```
model, tlh, tah, vlh, vah = fit()
```

In [14]:


```
plt.plot(tlh, label="train loss")
plt.plot(vlh, label="validation loss")
plt.legend()
plt.show()
```





In [15]:

```
plt.plot(tah, label="train accuracy")
plt.plot(vah, label="validation accuracy")
plt.legend()
plt.show()
```




In [16]:

```
hl_sizes=[32,64,128,256,512]
best_vl=[]
for hl in hl_sizes:
    _, tlh, tah, vlh, vah = fit(h=hl, epochs=20, lr=1e-3)
    best_vl.append(np.min(vlh))
```

In [17]:

```
plt.plot(hl_sizes, best_vl)
plt.show()
```



In [18]:

```
# h=512 gives best results
```


In [19]:

```
lr_vals=[10**(i-5) for i in range(5)]
best_vl=[]
for lr in lr_vals:
    _, tlh, tah, vlh, vah = fit(epochs=20, lr=lr)
    best_vl.append(np.min(vlh))
```

In [20]:

```
plt.plot(np.log10(lr_vals), best_vl)
```

```
plt.xlabel = "Learning rate as a power of 10"  
plt.show()
```



```
In [21]:
```

```
# lr=0.1 gives best results
```

```
In [22]:
```

```
full_test_loader = tud.DataLoader(test_set, batch_size=1500)  
for pts in full_test_loader:  
    test_pts=pts
```

```
In [23]:
```

```
best_m, tlh, tah, vlh, vah = fit(h=512, epochs=20, criterion=crit_bce, lr=0.1)
```

```
In [24]:
```

```
# mean loss  
label_pred = best_m(test_pts[0].reshape(-1,784).float())  
crit_bce(label_pred, test_pts[1])/64
```

```
Out[24]:
```

```
tensor(0.0010, grad_fn=<DivBackward0>)
```

```
In [25]:
```

```
label_pred=torch.argmax(label_pred, axis=1)
```

```
In [26]:
```

```
# accuracy  
print(torch.sum(label_pred==test_pts[1])/(label_pred.shape[0]))  
  
tensor(0.9810)
```

```
In [27]:
```

```
# very good test accuracy
```

```
In [31]:
```

```
class my_nn_3(tnn.Module):  
    def __init__(self, h=128, d_in=784, d_out=10):  
        super(my_nn_3, self).__init__()  
        self.linear1 = tnn.Linear(d_in, h)  
        self.dropout = tnn.Dropout(0.3)  
        self.linear2 = tnn.Linear(h, d_out)  
  
    def forward(self, x):  
        h_relu = self.linear1(x).clamp(min=0)  
        h_drop = self.dropout(h_relu)  
        y_pred = self.linear2(h_drop)
```

```
return y_pred
```

In [32]:

```
def fit_2(h=512, epochs=20, criterion=tnn.CrossEntropyLoss(), lr=1e-1):
    model = my_nn_3(h)
    opt = torch.optim.SGD(model.parameters(), lr=lr)
    tlh=[]
    tah=[]
    vlh=[]
    vah=[]
    vl_min=10**10
    for ep in range(epochs):
        model.train()
        training_loss=0
        acc=0
        for (idx, b) in enumerate(train_loader):
            xb=b[0].reshape(-1,784).float()
            yb=b[1]
            yb_pred = model(xb)
            loss=criterion(yb_pred, yb)
            training_loss+=loss
            acc+=torch.sum(torch.argmax(yb_pred, axis=1)==yb)
            opt.zero_grad()
            loss.backward()
            opt.step()
        tah.append(acc/50000)
        tlh.append(training_loss/50000)
        model.eval()
        with torch.no_grad():
            val_loss=0
            acc=0
            for (idx, b) in enumerate(val_loader):
                xb=b[0].reshape(-1,784).float()
                yb=b[1]
                yb_pred = model(xb)
                loss=criterion(yb_pred, yb)
                acc+=torch.sum(torch.argmax(yb_pred, axis=1)==yb)
                val_loss+=loss
            vah.append(acc/10000)
            vl=val_loss/64
            if vl<vl_min:
                best_model=dcopy(model)
                vl_min=vl
            vlh.append(vl)
    return best_model, tlh, tah, vlh, vah
```

In [33]:

```
best_m, _, _, _ = fit_2()
```

In [35]:

```
best_m
```

Out[35]:

```
my_nn_3(
  (linear1): Linear(in_features=784, out_features=512, bias=True)
  (dropout): Dropout(p=0.3, inplace=False)
  (linear2): Linear(in_features=512, out_features=10, bias=True)
)
```

In [37]:

```
label_pred = best_m(test_pts[0].reshape(-1,784).float())
print(crit_bce(label_pred, test_pts[1])/64)
label_pred=torch.argmax(label_pred, axis=1)
print(torch.sum(label_pred==test_pts[1])/(label_pred.shape[0]))
```

```
tensor(0.0009, grad_fn=<DivBackward0>)
tensor(0.9810)
```

In []:

```
# very similar results
```